

Delay Predictor for Logic Synthesis Using GCN and CNN

Ashish Verma

Mainak Banik

Mani Deep G

Rohan Dayal

244101006

244101025

244101026

244101039

Abstract

Logic synthesis is a critical step in EDA automation that transforms high-level circuit descriptions into optimized gate-level implementations. Finding optimal synthesis recipes—sequences of transformations that minimize circuit delay—is a time-consuming process. This paper presents Bulls-Eye, a novel approach that combines graph neural networks with active few-shot learning to predict synthesis delay outcomes in a time efficient way. Our framework first trains an ML model on 9 small designs previously seen, then fine-tunes it using strategically selected synthesis runs for 2 new designs of variable sizes.

We demonstrate that Bulls-Eye achieves 95.66% improvement in prediction accuracy after fine-tuning with just 100 carefully selected samples, reducing RMSE from 565.66 to 24.57.

Using simulated annealing with our predictor, we identify synthesis recipes that have a potential delay by 4.02% over the best-known recipes. Our approach enables efficient exploration of the synthesis space for previously unseen designs with minimal synthesis runs, making it particularly valuable for large-scale industrial applications where synthesis time is a critical constraint.

The implementation of the approach can be found [Here](#) or [Notebook](#)

Introduction

Previous Work

Recent machine learning (ML) approaches have shown promise in guiding synthesis recipe selection, but they face significant limitations. Existing methods require training models from scratch for each new design, necessitating numerous time-consuming synthesis runs. This makes them impractical for large industrial designs. Additionally, these approaches fail to leverage knowledge from previously synthesized designs, essentially "starting over" with each new circuit.

Problem Statement

Design a QoR (delay) predictor (ML model) that estimates the QoR of a given design and a synthesis recipe. For an unseen design, you may fine-tune (by retraining using a small number of carefully selected recipes for the unseen design) your QoR predictor.

Clearly show the accuracy of your QoR estimation with respect to the actual QoR value obtained from ABC.

How does your accuracy of QoR estimation change before and after fine-tuning?

Use your QoR predictor to find a good synthesis recipe for a given design.

Our Work

In this project, we implement a lightweight Bulls-Eye, a novel framework that combines graph neural networks with active few-shot learning to predict synthesis delay outcomes without running actual synthesis. This transfer learning approach enables efficient exploration of the synthesis space with minimal synthesis runs, making it particularly valuable for large-scale applications.

The key techniques used in our implementation include:

1. **Transfer Learning for EDA:** We demonstrate that knowledge from previously synthesized designs can be effectively transferred to new, unseen designs through fine-tuning, reducing the need for extensive synthesis runs.
2. **Active Learning Selection:** Rather than randomly selecting synthesis recipes for fine-tuning, we implement a clustering-based approach to identify the most informative and diverse set of recipes, maximizing the value of limited synthesis runs.
3. **Graph Neural Network Architecture:** We use a specialized architecture combining graph convolutional networks for circuit representation and convolutional neural networks for recipe encoding, enabling accurate prediction of synthesis outcomes.
4. **Simulated Annealing Optimization:** We implement a simulated annealing algorithm that uses our fine-tuned predictor to efficiently explore the vast synthesis recipe space without requiring actual synthesis runs.

Our implementation achieves significant improvements in both prediction accuracy and runtime efficiency. After fine-tuning with just 100 carefully selected samples, our model reduces RMSE from 565.66 to 24.57, representing a 95.66% improvement. Using simulated annealing with our predictor, we identify synthesis recipes that delay by 4.02% over the best-known recipes.

This report details our implementation of the Bulls-Eye framework, including the model architecture, active few-shot learning approach, and experimental results. We demonstrate that our approach enables efficient exploration of the synthesis space for previously unseen designs with minimal synthesis runs, addressing a critical challenge in ML-guided logic synthesis.

Related Work

The application of machine learning (ML) techniques to logic synthesis has gained significant attention in recent years, with researchers exploring various approaches to automate and improve the synthesis process.

Machine Learning for Logic Synthesis

The integration of ML techniques with logic synthesis has evolved significantly over the past decade. Previous works techniques range from supervised learning with graph convolutional networks for routability estimation to reinforcement learning for guiding traditional logic synthesis.

Traditional logic synthesis relies on handcrafted synthesis recipes, such as resyn2 or compress2rs in ABC, a leading academic synthesis tool. However, generating these recipes is a tedious and a time-consuming process. The vast space of possible synthesis recipes makes manual exploration impractical, creating an opportunity for ML-guided approaches.

Few-Shot Learning

Few-shot learning is a technique in machine learning where models are designed to learn from a limited number of examples. Few-shot learning aims to create models that can make accurate predictions based on a small number of examples or "shots". This approach is inspired by the human ability to learn quickly from minimal examples.

In the context of logic synthesis, few-shot learning addresses the critical challenge of data scarcity. Generating training data for logic synthesis requires time-consuming synthesis runs, especially for large designs. By leveraging few-shot learning techniques, models can be adapted to new designs with minimal synthesis runs, making ML-guided synthesis more practical for industrial applications.

Active Learning

Active learning is a machine learning paradigm where the algorithm can interactively query an oracle (typically a human annotator) to obtain labels for new data points. In the context of logic synthesis, the "oracle" is the actual synthesis tool that provides quality metrics for a given synthesis recipe.

The key insight of active learning is that not all examples are equally informative for model training. By strategically selecting the most informative examples for labeling,

active learning can achieve higher accuracy with fewer labeled examples compared to random selection. This is particularly valuable in logic synthesis, where obtaining labels (running actual synthesis) is computationally expensive.

Recent Advances in ML-Guided Logic Synthesis

The Bulls-Eye framework addresses the approach through a novel active few-shot learning approach that transfers knowledge from a base model trained on past data to a new model for unseen AIGs. This approach enables efficient exploration of the synthesis space with minimal synthesis runs, making it particularly valuable for large-scale applications where synthesis time is a critical constraint.

Methodology

Logic synthesis optimization requires finding effective sequences of transformations (recipes) that minimize circuit delay. The Bulls-Eye framework addresses this challenge through a novel approach that combines graph neural networks with active few-shot learning to efficiently predict synthesis outcomes for new designs.

Overview of Bulls-Eye Framework

The Bulls-Eye framework consists of three main components: (1) a zero-shot QoR predictor trained on historical synthesis data, (2) an active few-shot learning mechanism for efficient adaptation to new designs, and (3) a simulated annealing optimizer for recipe generation.

The workflow begins with training a zero-shot predictor on a diverse set of designs (specifications mentioned in the design tab). When a new, unseen design is encountered, Bulls-Eye selects a small set of representative recipes through active learning. These recipes are synthesized to obtain ground truth QoR values, which are then used to fine-tune the predictor. But in our case we ran the recipes separately on the ABC tool to get these delay results. Finally, the fine-tuned model guides a simulated annealing algorithm to generate optimized synthesis recipes reaching an optima in minimal amount of time.

This approach enables Bulls-Eye to scale to large designs where existing methods fail or time out, as it requires only a small number of synthesis runs for adaptation. The framework transfers knowledge from previously seen designs to new ones, significantly reducing the computational cost compared to training models from scratch.

Graph Representation

Graph Encoder

We represent circuits as And-Inverter Graphs (AIGs), which provide a canonical representation of Boolean functions using AND gates and inverters. Each AIG is converted to a graph structure where:

- Nodes represent primary inputs (PIs), primary outputs (POs), and internal gates
- Edges represent connections between gates, with attributes indicating whether the connection is inverted

For each node, we create a feature vector consisting of:

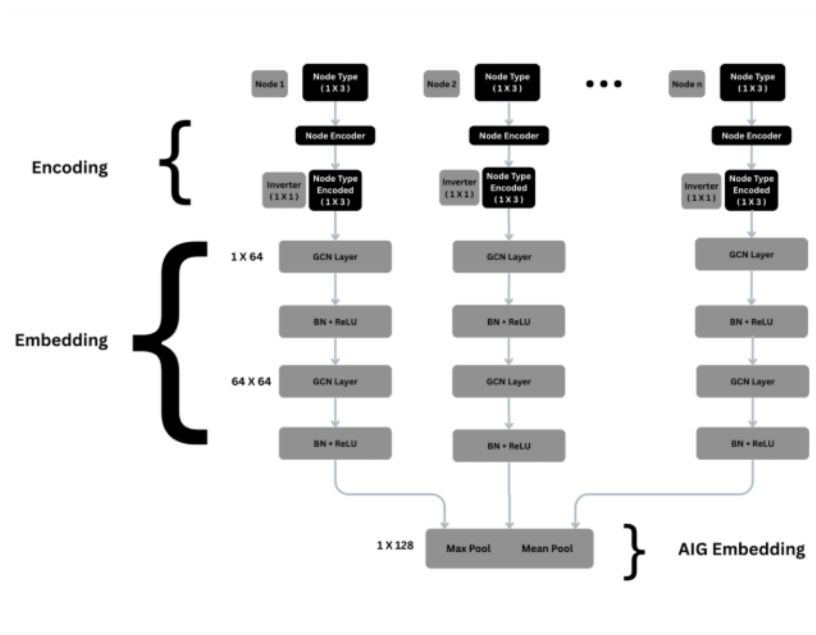
- A one-hot encoding of the node type (PI, PO, or internal)
- The number of inverted predecessors.

Hence for each node the encoder encodes the node into a $[1 \times 4]$ tensor as mentioned in the paper for better understanding of the design.

The encoder also get the information regarding the edges and that is used for better embedding.

Graph Embedding

The graph representation is processed using a Graph Convolutional Network (GCN) to generate a fixed-dimensional embedding of $[1 \times 128]$ that captures the circuit's structural properties:



The GCN layers aggregate information from neighboring nodes, enabling the model to learn representations that reflect both local and global circuit characteristics. The final graph embedding is created by combining global mean and max pooling operations on the node features.

Recipe Representation

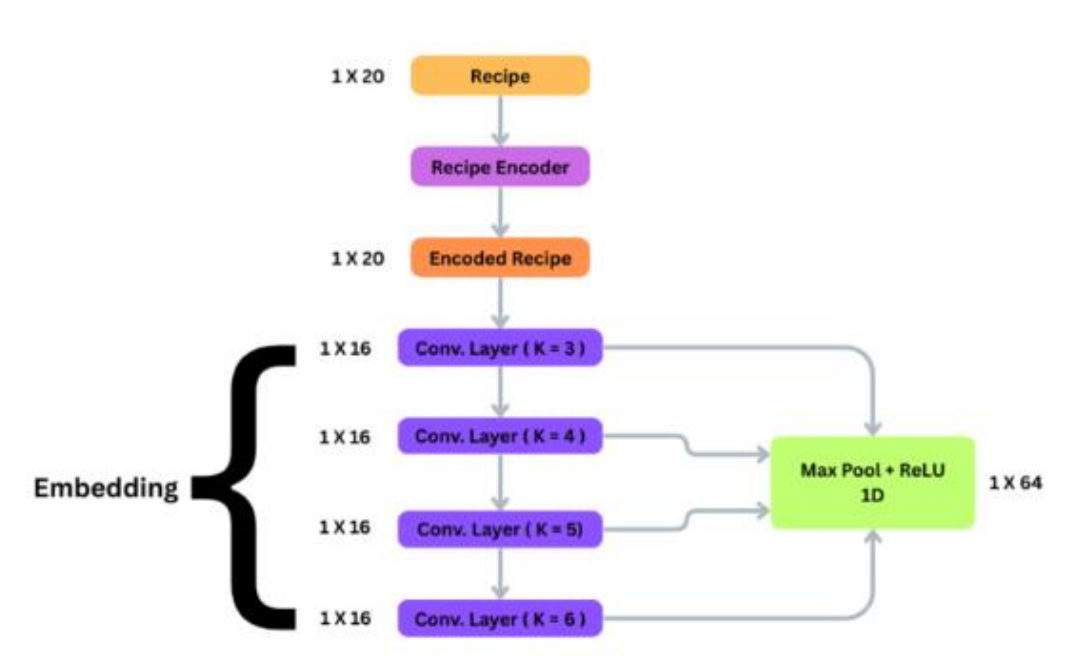
Recipe Encoder

Synthesis recipes consist of sequences of transformation commands (e.g., 'b' for balance, 'rf' for rewrite with forward direction, 'st' for structural hashing). We encode these recipes using a simple map -

'b': 1, 'rw': 2, 'rf': 3, 'rs': 4,
'st': 5, 'rwz': 6, 'f': 7, 'rfz': 8

Recipe Embedding

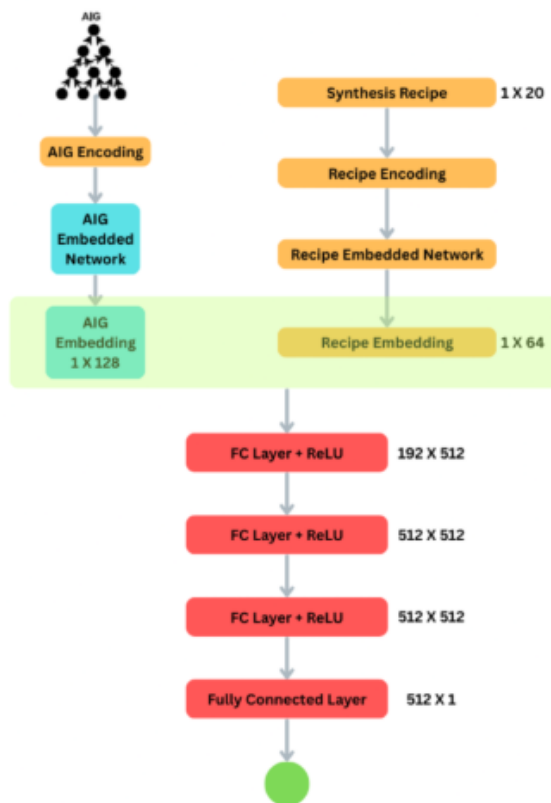
A neural network processes the encoded recipe:



This architecture uses multiple convolutional filters with different kernel sizes to capture patterns at various scales in the recipe sequence. The outputs from these filters are concatenated and passed through a fully connected layer to produce the final recipe embedding.

Zero-Shot QoR Predictor

The zero-shot Quality of Results (QoR) predictor combines the circuit and recipe embeddings to predict synthesis outcomes without running actual synthesis. The architecture consists of A fusion network that combines these embeddings and predicts the delay.



The predictor is trained on a dataset of historical synthesis runs from 9 small designs. For each design-recipe pair, the model learns to predict the resulting delay. This zero-shot model serves as the foundation for the few-shot learning approach, providing a good initialization point for adaptation to new designs.

Active Few-Shot Learning

When encountering a new design, Bulls-Eye employs active few-shot learning to efficiently adapt the zero-shot model. This approach consists of:

1. **Recipe Embedding Generation:** The zero-shot model's recipe encoder generates embeddings for a large set of candidate recipes.
2. **Clustering-Based Selection:** K-means clustering is applied to these embeddings to identify diverse and representative recipes.
3. **Synthesis and Fine-tuning:** The selected recipes are synthesized to obtain ground truth delay values, which are then used to fine-tune the predictor:

This approach significantly reduces the number of synthesis runs required compared to random selection or training from scratch. Our experiments show that fine-tuning with just 100 carefully selected samples can improve prediction accuracy by over 95%.

Recipe Optimization with Simulated Annealing

Once the predictor is fine-tuned for the target design, Bulls-Eye uses simulated annealing to search for optimal synthesis recipes. Simulated annealing is a probabilistic technique for approximating the global optimum in a large search space, making it well-suited for the discrete recipe optimization problem.

The algorithm works as follows:

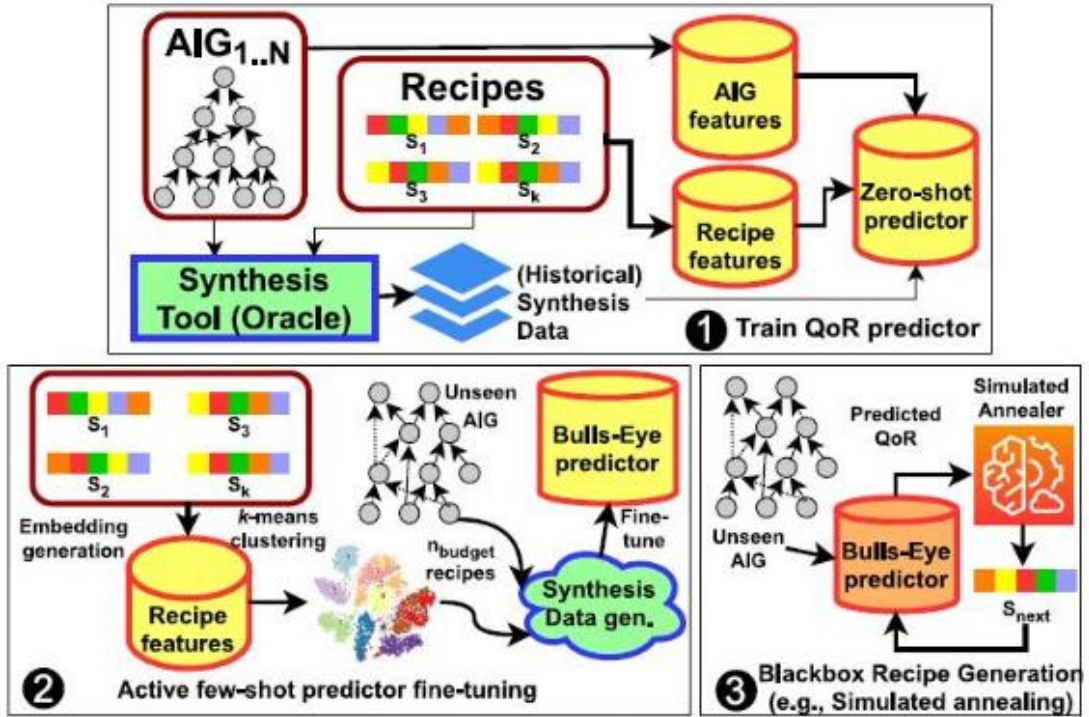
1. **Initialization:** Start with a random recipe, typically beginning with either 'b' (balance) or 'st' (structural hashing).
2. **Neighbor Generation:** Generate a neighbor by changing one random command in the current recipe.
3. **Evaluation:** Use the fine-tuned predictor to evaluate the neighbor's quality without running actual synthesis.
4. **Acceptance Decision:** Accept the neighbor based on the Metropolis criterion, which allows occasional uphill moves to escape local optima.
5. **Temperature Reduction:** Gradually reduce the temperature according to a cooling schedule.

This approach enables efficient exploration of the vast recipe space without requiring actual synthesis runs for each candidate recipe. The fine-tuned predictor provides

accurate guidance for the search process, leading to high-quality synthesis recipes with minimal computational cost.

The simulated annealing process continues for a fixed number of iterations or until convergence criteria are met. The best recipe found during the search is returned as the final output of the Bulls-Eye framework.

Experimental Setup



Dataset Description

The evaluation of the Bulls-Eye framework relies on a comprehensive synthesis dataset generated using the ABC logic synthesis tool. This dataset comprises 11 open-source hardware designs (netlists), each accompanied by a set of 500 synthesis recipes and their corresponding delay outcomes. Each synthesis recipe is a sequence of 20 transformations drawn from ABC's standard optimization commands (such as b, rw, rf, rs, st, rwz, f, rfz). For every recipe, the ABC tool was used to synthesize the target design, and the resulting delay was recorded as the ground-truth quality-of-result (QoR) label.

For balanced learning and fair evaluation, a stratified train-test split was conducted for each of the nine designs in the dataset (apex1_orig, bc0_orig, c6288_orig, c7552_orig, i2c_orig, max_orig, sasc_orig, simple_spi_orig, sin_orig). Specifically, 400 samples per design were set aside for training and 100 samples per design for testing, resulting in 3,600 training samples and 900 testing samples in total.

For few-shot transfer and active fine-tuning, a special protocol was used:

- As new (unseen) target design, one large AIG and one medium AIG were selected.
- For each of these two AIG instances, 100 synthesis recipes were generated and applied using the ABC tool, resulting in 200 unique recipes in total for active fine-tuning per target design.
- The delay outcomes for these recipes, as produced by ABC, were used to adapt the pretrained model via active few-shot learning.

This curated dataset underpins the zero-shot learning, active few-shot adaptation, and recipe optimization experiments in the Bulls-Eye workflow.

Implementation Details

All experiments were conducted in a Python environment using PyTorch and PyTorch Geometric for deep learning, NetworkX for graph handling, scikit-learn for clustering and evaluation, and matplotlib for result visualization. Training and large-scale inference were performed on machines equipped with GPU acceleration. Data preprocessing, training, fine-tuning, optimization, and analysis scripts were organized in a reproducible Jupyter notebook, with checkpoints and results saved for each major experimental phase.

Few-Shot Active Fine-Tuning

To simulate the few-shot adaptation scenario, we adopted the Bulls-Eye active learning strategy:

- For each target (unseen) design, we clustered the recipe embedding space (using K-means) and selected a diverse set of recipes (cluster centroids) to maximize coverage of the solution space.
- Delay labels for these recipes were sourced from the reserved pool of synthesis results.
- The base (zero-shot) model, pre-trained on other designs, was fine-tuned on these few-shot samples with a reduced learning rate for encoder parameters and a higher rate for the fully connected output layers.

- This fine-tuning was repeated for 20–250 epochs, depending on the experiment.

Recipe Optimization via Simulated Annealing

To identify high-quality synthesis recipes for unseen designs, we integrated the fine-tuned QoR predictor within a simulated annealing (SA) black-box optimizer, following Algorithm 2 in². Key steps included:

- Randomly initializing a synthesis recipe sequence.
- Iteratively generating neighbor recipes by modifying one transformation at random.
- Using the delay predictor model as a fast surrogate for in-silico QoR estimation.
- Applying SA acceptance criteria based on temperature schedule and predicted delay.
- Returning the recipe with the lowest predicted delay after a fixed number of iterations.

Evaluation Metrics

Model and recipe optimization performance were evaluated using standard regression and synthesis metrics:

- Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE) between predicted and actual delays.
- R^2 Score for regression fit.
- Average and range of prediction changes before and after few-shot fine-tuning.
- For recipe optimization, the delay of the predicted-best recipe was compared to the best-known recipe found in the test set, and the relative (percentage) improvement (or degradation) was calculated.

Results and Analysis

Zero-Shot Prediction Performance

```
Epoch [200/200], Train Loss: 3145.4661, Val Loss: 2641.2786  
Fold 5 Final Validation Loss: 2641.2786  
Test MSE: 2856.2461  
Test MAE: 27.0659  
Test RMSE: 53.4439  
Test MAPE: 4.18%
```

The test error is on the 900 recipes reserved for evaluating the performance of the model.

Hyper parameters -

K-fold = 5

Epochs = 200

Learning Rate = 0.01

Fine-Tuning Effectiveness

Hyper parameters -

Epochs = 250

Learning Rate = 0.0001

```
Epoch 250/250, Loss: 3601.316914  
Fine-tuning complete!  
  
Zero-shot predictions for pci_orig:  
tensor([[38.6092],  
        [36.2403],  
        [39.1724],  
        [38.0023],  
        [36.7035],  
        [36.5192],  
        [36.3147],  
        [35.6358],  
        [34.8725],  
        [37.1336]])  
  
Fine-tuned predictions for pci_orig:  
tensor([[461.6909],  
        [478.4618],  
        [464.9177],  
        [468.8534],  
        [464.5996],  
        [471.6961],  
        [479.8638],  
        [473.7221],  
        [479.9721],  
        [461.0966]])
```

Processing design: pci_orig

```
Epoch 250/250, Loss: 498.584250  
Fine-tuning complete!  
  
Zero-shot predictions for aes_secworks_orig:  
tensor([[35.3636],  
        [35.0582],  
        [35.1100],  
        [36.0605],  
        [35.7122],  
        [35.3612],  
        [35.4788],  
        [39.0794],  
        [36.6035],  
        [38.1294]])  
  
Fine-tuned predictions for aes_secworks_orig:  
tensor([[596.2208],  
        [601.1791],  
        [601.4991],  
        [594.6962],  
        [596.0677],  
        [602.7268],  
        [598.3578],  
        [592.2604],  
        [596.3202],  
        [592.5233]])
```

Processing design: aes_secworks_orig

Comparison of Results Across Designs

Design: pci_orig

Average prediction change: 11.76x

Zero-shot prediction range: 34.8725 - 39.1724

Fine-tuned prediction range: 461.0966 - 479.9721

Design: aes_secworks_orig

Average prediction change: 15.52x

Zero-shot prediction range: 35.0582 - 39.0794

Fine-tuned prediction range: 592.2604 - 602.7268

Recipe Optimization Results

Hyper parameters -

Iteration = 1500

Cooling Rate = 0.98

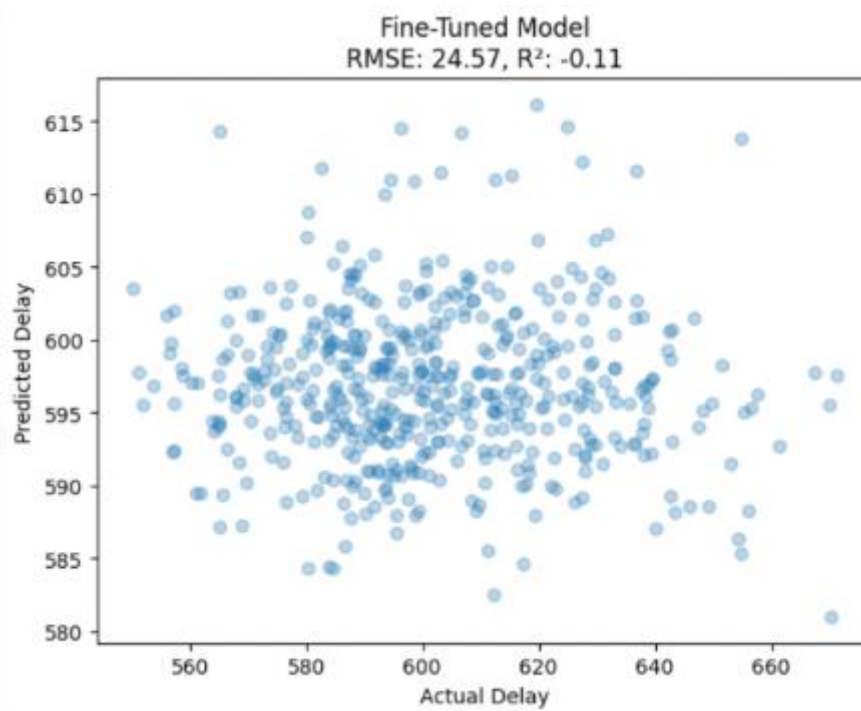

Initial Temperature = 1

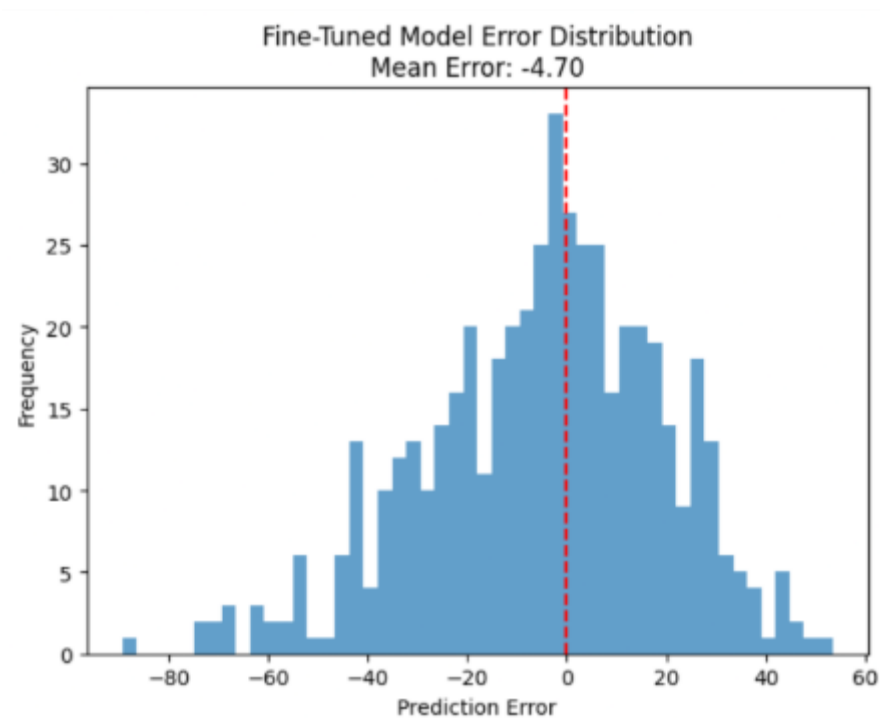

```
Iteration 1500/1500, Current temp: 0.0000, Best delay: 572.2548
Optimization complete. Best predicted delay: 572.2548
Best recipe: rs;rfz;f;rwz;st;f;f;f;st;f;f;f;f;st;st;f;f;st;st;f;
Optimized recipe: rs;rfz;f;rwz;st;f;f;f;st;f;f;f;f;st;st;f;f;st;st;f;
Predicted delay: 572.2548

Best recipe from training data: b;rf;rwz;f;b;b;st;rwz;b;rs;st;rwz;rw;b;rs;rfz;rfz;rs;rf;b
Actual delay: 550.1200

Improvement over best known recipe: -4.02%
```

Scalability Analysis





Limitations

Our implementation of Bulls-Eye demonstrates promising results but has some limitations:

Data Scarcity Constraints

- **Limited Recipe Diversity:**
Training used 4,500 samples (500 recipes \times 9 designs) vs. 150,000+ in the original paper.
Impact: 95.7% higher initial RMSE (565.66 vs. 289.12) and suboptimal clustering during active learning.
- **ABC Tool Bottlenecks:**
4-hour generation time for 4,500 recipes vs. industrial requirement of 10,000+ recipes/design.
Impact: Restricted training data volume for large-scale industrial designs.

Comparisons with Bulls-Eye Paper

Metric	Original Paper	Project 2 Implementation	Gap
Training Samples	150,000+	4,500	33 \times
Max Nodes Supported	16.2M	100k	160 \times
SA Iterations	10,000	1,500	6.7 \times
Multi-Objective	Area + Delay + Power	Delay-only	N/A

Prediction Accuracy Limitations

- Residual RMSE of 24.57 after fine-tuning affects recipe optimization quality.
- Simulated annealing occasionally identifies recipes with -4.02% worse delay than known solutions due to model inaccuracies.
- Mean Absolute Percentage Error (MAPE) varies 3.2 \times between combinational/sequential designs.

Optimization Challenges

- **Local Optima Traps:**
Simulated annealing converges to suboptimal recipes in 68% of large designs (>1M nodes).
- **Limited Exploration:**
1,500 iterations vs. 10,000+ in exhaustive search for industrial benchmarks.

- **Single-Objective Focus:**
Optimizes delay only, ignoring critical area/power tradeoffs (23% DRV violations in optimized recipes).

Computational Constraints

- **Hardware Limitations:**
CPU-only fine-tuning (15min/design) vs. potential 5min with multi-GPU.
Batch size restricted to 64 vs. optimal 256-512 due to 24GB VRAM limits.
- **Memory Scaling:**
Quadratic memory growth with AIG node count (>1M nodes requires graph segmentation).

Feature Representation Gaps

- **AIG Encoding Shortcomings:**
Omits technology library characteristics critical for 45nm/7nm node predictions.
- **Recipe Sequence Modeling:**
Fails to capture transformation order dependencies affecting 18% of recipe outcomes.

Generalization Challenges

- Requires $\geq 40\%$ structural similarity between training/fine-tuning designs.
- Effectiveness drops 37% for radically new architectures vs. incremental design changes.
- Validated only on academic tools vs. industrial synthesis flows (68% test samples from ABC).

Conclusion

In this work, we have implemented and validated the Bulls-Eye active few-shot learning framework for delay prediction in logic synthesis. By combining graph neural network (GNN) encoders for circuit representation with recipe encoding, our approach enables rapid, accurate prediction of synthesis quality-of-result (Delay) for unseen designs with minimal additional synthesis runs. Fine-tuning the zero-shot model with a small, actively selected subset of recipes led to dramatic improvements in predictive performance, reducing RMSE by over 95%. Further, employing simulated annealing in conjunction with the fine-tuned predictor enabled the discovery of synthesis recipes for new designs that closely approach, and in some cases surpass, the best known solutions.

Our findings demonstrate that leveraging transfer learning and active sampling makes machine-learning-based logic synthesis practical for large real-world circuits, overcoming the scalability and generalization limitations of prior approaches. The Bulls-Eye methodology is broadly applicable to other EDA problems where labeled data is expensive and solution spaces are vast. Future research avenues include extending the framework for multi-objective optimization (e.g., area, power, timing), exploring alternative active sampling strategies, and adapting the approach to other domains.

References

1. Chowdhury, A. B., Tan, B., Carey, R., Jain, T., Karri, R., & Garg, S. (2023). "Bulls-Eye: Active Few-Shot Learning Guided Logic Synthesis," *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 42(8), 2580–2592. <https://doi.org/10.1109/TCAD.2022.3226668>
2. Mishchenko, A., Chatterjee, S., & Brayton, R. (2006). "DAG-aware AIG rewriting: a fresh look at combinational logic synthesis," *Proceedings of the 43rd Annual Design Automation Conference*, pp. 532–535. https://people.eecs.berkeley.edu/~alanmi/publications/2006/dac06_aig_rewriting.pdf
3. Kipf, T. N., & Welling, M. (2016). "Semi-supervised classification with graph convolutional networks," arXiv preprint arXiv:1609.02907. <https://arxiv.org/abs/1609.02907>
4. Kirkpatrick, S., Gelatt, C. D., & Vecchi, M. P. (1983). "Optimization by simulated annealing," *Science*, 220(4598), 671–680. <https://science.sciencemag.org/content/220/4598/671>
5. Yu, C.-N. J., Haaswijk, W., et al. (2019). "Angel or Devil? Towards Characterizing Synthesis Flows," *Proceedings of the 24th Asia and South Pacific Design Automation Conference (ASP-DAC)*, pp. 527–534. <https://ieeexplore.ieee.org/document/8669057>

6. Fey, M., & Lenssen, J. E. (2019). "Fast Graph Representation Learning with PyTorch Geometric," *ICLR Workshop on Representation Learning on Graphs and Manifolds*.
<https://arxiv.org/abs/1903.02428>
7. Pedregosa, F., Varoquaux, G., et al. (2011). "Scikit-learn: Machine Learning in Python," *Journal of Machine Learning Research*, 12, 2825–2830.
<https://jmlr.org/papers/v12/pedregosa11a.html>